

LEIGH HALLIDAY

Efficient Search in Rails with Postgres

Efficient Search in Rails with Postgres

Allowing your users to find the data they are looking for is very important. For example, it can mean the difference between a sale or no sale, in the case of e-commerce applications.

The type of data you have may help determine the best way to find it. Or, imagine another situation, where we are developing a stock market research website where the user can search for a company using a single text input.

In this eBook we show how we speed up a search query from seconds to milliseconds. We walk through using exact matches, similarity matches with trigrams, partial matches with ILIKE, and natural language matches.

ABOUT

They may enter:

- ▶ **AAPL:** The stock market symbol for the company.
- ▶ **Apple:** The (partial) name of the company.
- ▶ **Integrated hardware and software:** When they don't know exactly what they're looking for, and want to find the best match using natural language.

In this ebook, we will cover how to efficiently find the correct data for our user for all three scenarios mentioned above using only Ruby on Rails and Postgres. We will see when you might choose one over the other along with the pros and cons of each approach, aiming to optimize our queries along the way.

1	Postgres Feature	Typical Use Case	Can Be Indexed?	Performance
2	-----	-----	-----	-----
3	Exact \(\=\)	Exact match or user selecting from predefined list	Yes \(\BTREE\)	Great
4	LIKE/ILIKE	Wildcard\~style search for small data	Sometimes	Unpredictable
5	Trigram \(\pg_trgm\)	Similarity search for names, etc	Yes \(\GIN, GiST\)	Good
6	Full Text Search	Natural language search	Yes \(\GIN, GiST\)	Good

What Data Are We Using?

The data used in the following examples consists of 253k records in a table called `companies`. There are about 3k companies listed on the [NASDAQ](#) Stock Exchange, 700 from the [TSX](#) (Toronto Stock Exchange), and then an additional 250k records consisting of fake data generated using [Faker](#).

Our goal is for all three of the user's searches to return this record:

JSON

```
1  {
2    "id": 256651,
3    "exchange": "NASDAQ",
4    "symbol": "AAPL",
5    "name": "Apple Inc.",
6    "description": "Apple designs a wide variety of consumer electronic devices,
including smartphones (iPhone), tablets (iPad), PCs (Mac), smartwatches (Apple Watch),
and TV boxes (Apple TV), among others. The iPhone makes up the majority of Apple's
total revenue. In addition, Apple offers its customers a variety of services such as
Apple Music, iCloud, Apple Care, Apple TV+, Apple Arcade, Apple Card, and Apple Pay,
among others. Apple's products run internally developed software and semiconductors,
and the firm is well known for its integration of hardware, software and services."
7  }
```

If you would like to follow along, the [source code](#) is available, and by running `rails db:seed` it will generate the 253k records mentioned above.

Exact Matches with Equals

Exact matches are perfect for when the user knows exactly what they are searching for (a stock symbol, an email address or username, a referral code), or when they are choosing from a predefined list of options. This is the most restrictive approach to searching, but it is easy to optimize using a standard B-tree index in Postgres. No additional Ruby gems or Postgres extensions are required for it to work.

Pros:

- ▶ It is simple: Works out of the box
- ▶ It is performant: A standard B-tree index does wonders

Cons:

- ▶ It must be an exact match
- ▶ It is case sensitive

The easiest way to find a single matching company record in Rails is by using the `find_by` method.

RUBY

```
1 Company.find_by(symbol: 'AAPL')
```

The above Ruby code will produce the following SQL:

SQL

```
1 SELECT "companies".* FROM "companies" WHERE "companies"."symbol" = 'AAPL' LIMIT 1
```

Optimizing an Exact Match Query

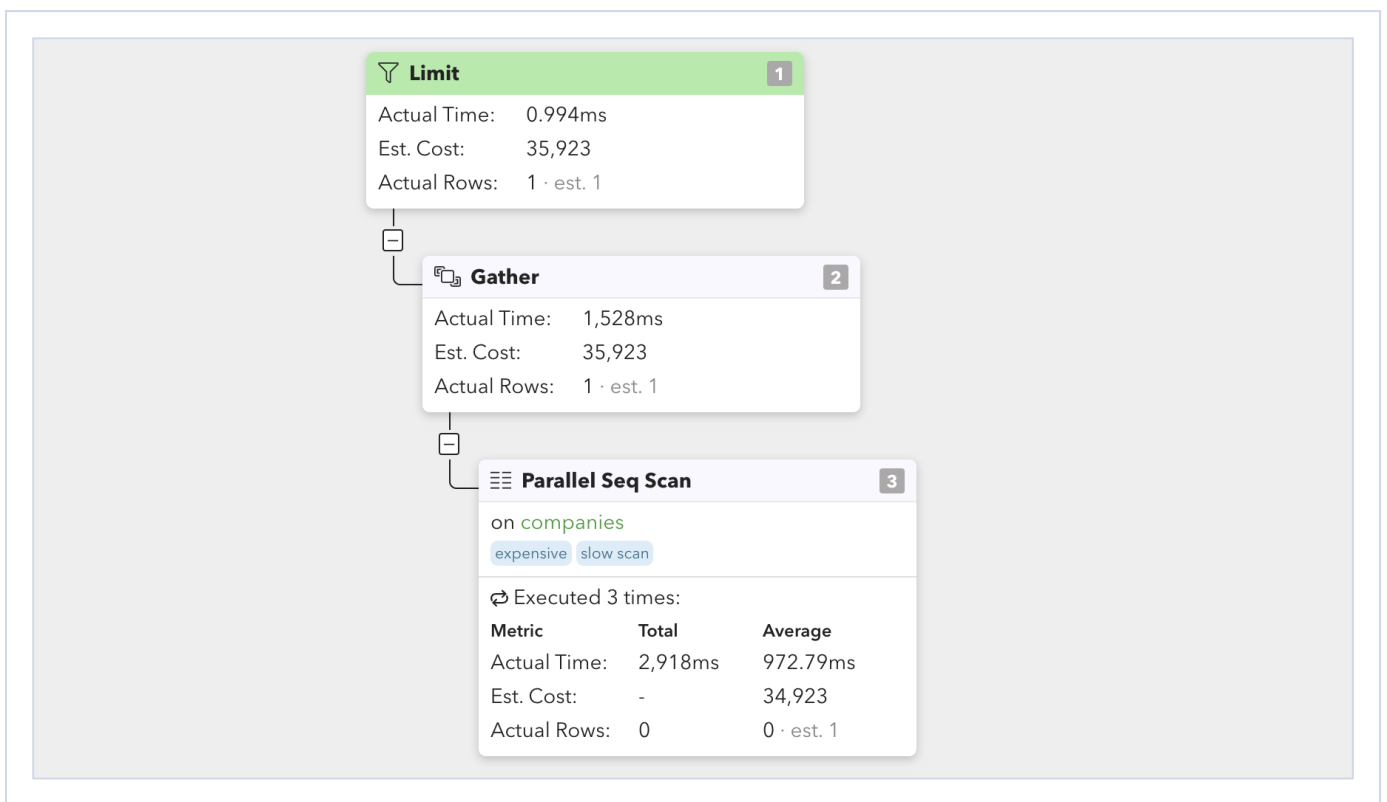
Because we knew that we would be searching for companies by their symbol, we created an index on this column at the same time that the table was created.

RUBY

```
1 class CreateCompanies < ActiveRecord::Migration[6.0]
2   def change
3     create_table :companies do |t|
4       t.string :exchange, null: false
5       t.string :symbol, null: false
6       t.string :name, null: false
7       t.text :description, null: false
8       t.timestamps
9
10      t.index :symbol
11      t.index %i[exchange symbol], unique: true
12    end
13  end
14 end
```

By adding an index, it takes a query that would be close to 1500ms without an index, to just 7ms. That's an improvement of **over 200x!**

Before (Sequential Scan):



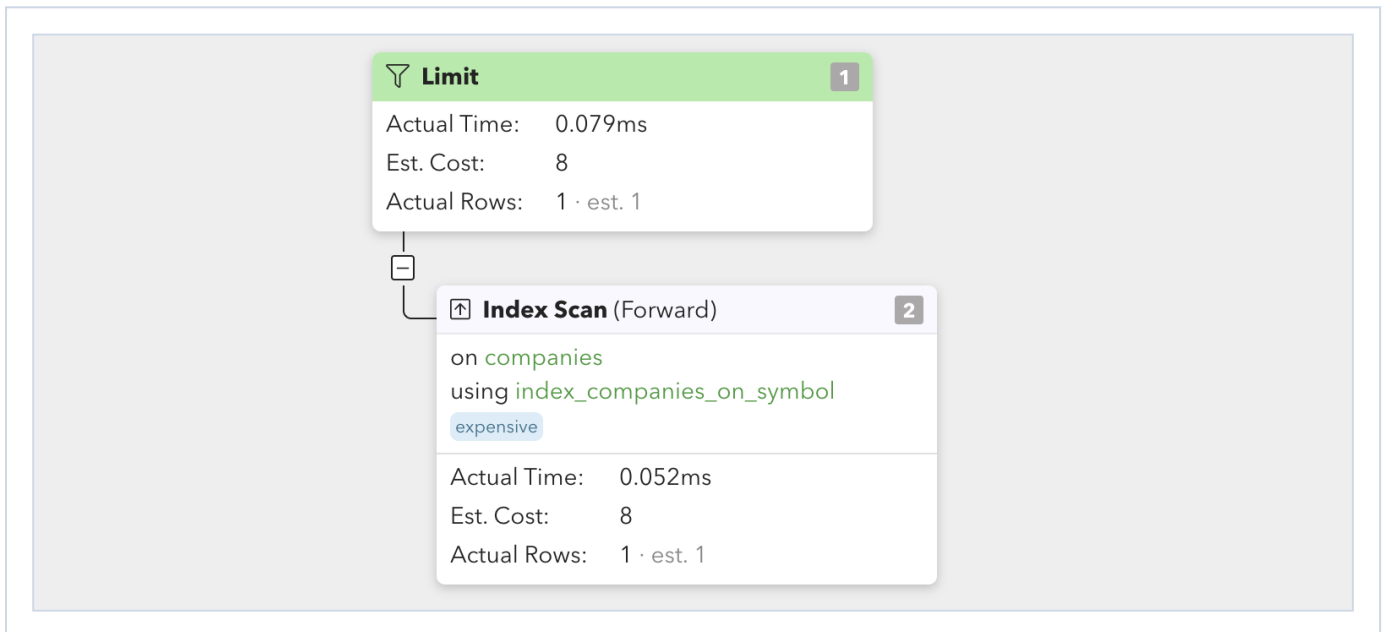
HOW COMPANIES ARE USING PGANALYZE

ATLASSIAN

Case Study: How Atlassian and pganalyze are optimizing Postgres query performance

[Read The Story](#)

After (Index Scan):



Partial / Case Insensitive Matches with ILIKE

Exact matches can be extremely restrictive because they only help in a limited set of search use cases.

An alternative is to provide wildcard (partial) searching using **LIKE** and its case insensitive partner **ILIKE**.

Pros:

- ▶ The user doesn't have to worry about the case matching
- ▶ The user can enter partial matches
- ▶ No additional Ruby gems required

Cons:

- ▶ Index usage is unpredictable
- ▶ Spelling must be accurate (Applo would not find Apple)

To find Apple from our companies table, the following two queries will both do the trick. The first will simply use **ILIKE** to allow for case insensitive searching, while the second query will use the **%** wildcard before and after the term, finding any company that has **apple** in its name.

RUBY

```
1 Company.where('companies.name ILIKE ?', 'apple inc.').take
2 Company.where('companies.name ILIKE ?', '%apple%').take
```

SQL

```
1 -- case insensitive
2 SELECT companies.*
3   FROM companies
4  WHERE companies.name ILIKE 'apple inc.'
5   LIMIT 1;
6
7 -- wildcard search
8 SELECT companies.* FROM companies WHERE companies.name ILIKE '%apple%' LIMIT 1;
```

Using Indexes with LIKE and ILIKE

The `pg_trgm` extension allows `GIN` and `GIST` indexes on text columns, which can be used to speed up `LIKE` and `ILIKE` searches. That said, even after indexing the `companies.name` column this way, the queries in our example still don't use this index and instead continue doing a full table sequence scan, providing no performance increase. This is due to the structure of the associated trigrams, and the Postgres planner determining a sequential scan is still more effective.

`GIN/GIST` indexes together with `pg_trgm` can sometimes be used for `LIKE` and `ILIKE`, but query performance is unpredictable when user-generated input is presented.

To add a `GIN` trigram index to the name column, we will first need to enable the `pg_trgm` extension. Note the `pg_trgm` extension ships with a standard Postgres installation and is available with major cloud providers, making it safe to rely on it:

RUBY

```
1 class EnableTrigramExtension < ActiveRecord::Migration[6.0]
2   def change
3     enable_extension :pg_trgm
4   end
5 end
```

With the extension enabled, we can add an index to our `name` column:

RUBY

```
1 class AddTrigramIndexCompaniesName < ActiveRecord::Migration[6.0]
2   disable_ddl_transaction!
3
4   def change
5     add_index :companies,
6             :name,
7             opclass: :gin_trgm_ops,
8             using: :gin,
9             algorithm: :concurrently,
10            name: 'index_companies_on_name_trgm'
11   end
12 end
```

On the plus side, this is the **very same** extension and index we will be using in the following section on similarity matches, so all is not lost!

Similarity Matches with Trigrams

You searched **"Applo"**... did you mean **"Apple"**?

Trigrams allow us to find matches based on similarity. This means that even if the user misspells a company name, as long as their spelling is similar to the real one, we'll be able to find the correct match.

Pros:

- ▶ Misspellings are no problem matching
- ▶ Searches are case insensitive
- ▶ Searches can be indexed

Cons:

- ▶ Does not replace the need for natural language search

What is a Trigram?

Trigrams involve breaking a string into groups of three consecutive letters. For the word **Apple**, it would be broken into **three** groups: app, ppl, ple.

This process of breaking a piece of text into smaller groups allows you to compare the groups of one word to the groups of another word. Knowing how many groups are shared between the two words allows you to make a comparison between them, based on how similar their groups are.

Let's see an example of this functionality in SQL:

SQL

```
1 SELECT
2   show_trgm('Apple'), -- {" a"," ap",app,"le ",ple,ppl}
3   show_trgm('Applo'), -- {" a"," ap",app,"lo ",plo,ppl}
4   similarity('Apple', 'Apple'), -- 1
5   similarity('Applo', 'Apple'), -- 0.5
6   similarity('Oapple', 'Apple'), -- 0.33
7   'Applo' % 'Apple' -- TRUE
```

The `show_trgm` function takes a string and returns us a list of trigrams. You'll notice something interesting here, that there are more than three results... there are actually six! Two spaces were added to the beginning of the string, and a single space was added to the end.

The first reason is that it allows similarity calculations on words with less than three characters, such as `Hi`.

Secondly, it ensures the first and last characters are not overly de-emphasized for comparisons. If we used only strict triplets, the first and last letters in longer words would each occur in only a single group: with padding they occur in three (for the first letter) and two (for the last). The last letter is less important for matching, which means that `Apple` and `Applo` are more similar than `Apple` and `Oapple`, even though they are both off by a single character.

The `similarity` function provides us a percentage of shared trigrams between zero and one. Zero meaning that they aren't similar at all, one meaning they are a perfect match.

Lastly, we have the `%` operator, which gives you a boolean of whether two strings are similar. By default, Postgres uses the number 0.3 when making this decision, but you can always [update this setting](#).

Using Trigrams in Rails

The `pg_trgm` extension must be enabled to use Trigram functionality. We have already seen this when adding an index to improve `ILIKE` search, but if you skipped forward to this section:

RUBY

```
1 class EnableTrigramExtension < ActiveRecord::Migration[6.0]
2   def change
3     enable_extension :pg_trgm
4   end
5 end
```

Typically we not only care whether a match is considered similar or not, but also **how** similar it is. We want more similar results before less similar

ones. To accomplish this, let's create a `scope` that will not only filter results using the `%` (similarity) operator, but will also sort them by how similar they are, from most to least (descending order).

RUBY

```
1 class Company < ApplicationRecord
2   scope :name_similar,
3     lambda { |name|
4       quoted_name = ActiveRecord::Base.connection.quote_string(name)
5       where('companies.name % :name', name: name).order(
6         Arel.sql("similarity(companies.name, '#{quoted_name}') DESC")
7     )
8   }
9 end
10
11 Company.name_similar('Applo').first
```

The SQL produced with the above Ruby code looks like:

SQL

```
1  SELECT companies.*
2  FROM companies
3  WHERE companies.name % 'Applo'
4  ORDER BY similarity(companies.name, 'Applo') DESC
5  LIMIT 1;
```

Optimizing Trigram Search

Trigram (similarity) search can be optimized by adding either a GIN or GiST index to your column, using a special `opclass` of `gin_trgm_ops`. We will be working with GIN indexes because Postgres recommends them when [working with text](#):

GIN indexes are the preferred text search index type. As inverted indexes, they contain an index entry for each word (lexeme), with a compressed list of matching locations. Multi-word searches can find the first match, then use the index to remove rows that are lacking additional words. GIN indexes store only the words (lexemes) of tsvector values, and not their weight labels. Thus a table row recheck is needed when using a query that involves weights.

GIN INDEXES

In Rails, the type of index and opclass can be set using options on the `add_index` method. Note that we also included `disable_ddl_transaction!` at the top of this migration, and set the algorithm to `concurrently` so that this index will be added [concurrently](#). This is useful (or crucial) when you have large amounts of rows and want to avoid locking the table while the

index is being applied.

RUBY

```
1 class AddTrigramIndexCompaniesName < ActiveRecord::Migration[6.0]
2   disable_ddl_transaction!
3
4   def change
5     add_index :companies,
6               :name,
7               opclass: :gin_trgm_ops,
8               using: :gin,
9               algorithm: :concurrently,
10              name: 'index_companies_on_name_trgm'
11   end
12 end
```

Natural Language Matches with Full Text Search

Full Text Search allows a user to search through large documents by simply describing in English words what they want. In our example, searching „integrated hardware and software“ allows us to find Apple based on their description, even though those words don't show up in exactly the same form or order.

Pros:

- ▶ Users can use natural language
- ▶ Easily search through large documents
- ▶ Searches are very performant

Cons:

- ▶ Can be difficult to configure

Minimal Full Text Search in Rails Setup

We will be using a gem called `pg_search` to make Postgres' Full Text Search easier to integrate into our Rails application.

After installing the `pg_search` gem, it's extremely easy to add Full Text Search to your Rails model:

RUBY

```
1 class Company < ApplicationRecord
2   include PgSearch::Model
3   pg_search_scope :search, against: :description
4 end
```

Then we can call `search` and search through the companies' description column using a natural language search.

RUBY

```
1 Company.search('integrated hardware and software').first
```

Unfortunately, there are a **number of issues** with the simple setup above. It is slow (15 seconds), it isn't aware that we are searching in the English language, and what if we wanted to search on both the `name` and the `description` columns, giving precedence to the `name`? We will learn how to solve all of these issues below.

Interested in learning more about searching in Rails with Postgres? Check out our blog posts!

[Similarity in Postgres and Rails using Trigrams \(click to read article\)](#)

[Full Text Search in Milliseconds with Rails \(click to read\)](#)

PGANALYZE

The Foundations of Full Text Search

By taking time to understand the underlying concepts of Postgres' Full Text Search, we can better solve the issues that have arisen in our simple search configuration.

SQL

```
1 SELECT
2   to_tsvector('english', 'the firm is well known for its integration of hardware,
3   -- 'firm':2 'hardwar':10 'integr':8 'known':5 'servic':13 'softwar':11 'well':4
4   software and services');
```

In the above SQL we have some text (taken from the Apple company description); often referred to as a **document** when talking about Full Text Search. A document must be parsed and converted into a special data type called a **tsvector**, which we did using the function **to_tsvector**.

The **tsvector** data type is comprised of **lexemes**. Lexemes are normalized key words which were contained in the document that will be used when searching through it. In this case we used the English language dictionary to normalize the words, breaking them down to their root. This means that

„integration“ became „integr“, „services“ became „servic“, with **very common words** being removed completely, to avoid false positives.

SQL

```
1 SELECT to_tsvector(  
2     'english',  
3     'the firm is well known for its integration of hardware, software and  
4     services'  
5 )  
6 @@ to_tsquery('english', 'integrated & hardware & and & software'), -- TRUE  
7 to_tsquery('english', 'integrated & hardware & and & software'); -- 'integr' &  
8 'hardwar' & 'softwar'
```

The `@@` operator allows us to check if a query (data type `tsquery`) exists within a document (data type `tsvector`). Much like `tsvector`, `tsquery` is also normalized prior to searching the document for matches. Although `pg_search` handles this for us, search terms passed to `to_tsquery` must be separated by `&`.

SQL

```
1 SELECT ts_rank(  
2     to_tsvector(  
3     'english',  
4     'the firm is well known for its integration of hardware, software and  
5     services'  
6 ),  
7 to_tsquery('english', 'integrated & hardware & and & software')  
8 );  
9 -- 0.2669
```

The `ts_rank` function takes a `tsvector` and a `tsquery`, returning a number that can be used when sorting the matching records, allowing us to sort the results from highest to lowest ranking.

Searching a Single Column

Combining these concepts together, we can find companies with a matching query string:

SQL

```
1  SELECT id,  
2      exchange,  
3      symbol,  
4      name,  
5      ts_rank(  
6          to_tsvector('english', description),  
7          to_tsquery('english', 'integrated & hardware & and & software')  
8      ) AS rank  
9  FROM companies  
10 WHERE to_tsvector('english', description)  
11      @@ to_tsquery('english', 'integrated & hardware & and & software')  
12 ORDER BY rank DESC  
13 LIMIT 2;
```

This will return Apple as the number one rank, followed by a company called Vecima Networks Inc. Interestingly, if the language `english` isn't specified (defaulting to a dictionary called `simple`), these companies will come back in the opposite order.

JSON

```
1  [  
2    {  
3      "id": 9,  
4      "exchange": "NASDAQ",  
5      "symbol": "AAPL",  
6      "name": "Apple Inc.",  
7      "rank": 0.29986802  
8    },  
9    {  
10     "id": 3640,  
11     "exchange": "TSX",  
12     "symbol": "VCM",  
13     "name": "Vecima Networks Inc.",  
14     "rank": 0.2642635  
15   }  
16 ]
```

Searching Multiple Columns

Up until this point we have only seen a single column being searched at a time, but what if we want our Full Text Search to look at both the `name` **AND** the `description` columns? And what if we wanted to give precedence to the `name` column?

Data types of `tsvector` can be concatenated with the `||` operator, and precedence (weight) can be given to a certain column using the `setweight` function. Valid weighting options are: A, B, C or D.

SQL

```
1  SELECT id,
2      exchange,
3      symbol,
4      name,
5      ts_rank(
6          setweight(to_tsvector('english', name), 'A')
7          || setweight(to_tsvector('english', description), 'B'),
8          to_tsquery('english', 'integrated & hardware & and & software')
9      ) AS rank
10 FROM companies
11 WHERE setweight(to_tsvector('english', name), 'A')
12        || setweight(to_tsvector('english', description), 'B')
13        @@ to_tsquery('english', 'integrated & hardware & and & software')
14 ORDER BY rank DESC
15 LIMIT 2;
```

Configuring pg_search

With our knowledge of `tsvector` concatenation, using `setweight` to specify column precedence, and setting `english` as the language of our choice, let's configure `pg_search` to build this query for us.

RUBY

```
1  class Company < ApplicationRecord
2      include PgSearch::Model
3
4      pg_search_scope :search,
5          against: { name: 'A', description: 'B' },
6          using: { tsearch: { dictionary: 'english' } }
7  end
```


Which can be queried with:

RUBY

```
1 Company.search('integrated hardware and software').first
```

Optimizing Full Text Search

We still have a major issue on our hands! The query takes **15 seconds** to complete, **much** too long for it to be used in a production system. Why does it take so long?

It is slow for two reasons:

- ▶ The tsvectors are being calculated for the entirety of our data on every query
- ▶ No index is being used

As we are venturing into the territory of more custom Postgres functionality, not easily supported by the Rails schema file in Ruby, we'll want to **switch the schema format** from `:ruby` to `:sql`. This line can be added to the application.rb file:

Which can be queried with:

RUBY

```
1 config.active_record.schema_format = :sql
```

To solve the first problem, we can create a migration to add a **generated** and **stored** column that will store the resulting **tsvector** value. Note that generated columns are only supported with Postgres 12 and newer, but there are alternatives for older Postgres versions. This generated column actually occupies space in our table and gets written on each INSERT/UPDATE. This also means that when we add a generated column to a table, it will require a rewrite of the table to actually set the values for all existing rows. This may take some time for a large table and block other operations on the database.

RUBY

```
1 class AddSearchableToCompanies < ActiveRecord::Migration[6.0]
2   def up
3     execute <<-SQL
4       ALTER TABLE companies
5         ADD COLUMN searchable tsvector GENERATED ALWAYS AS (
6           setweight(to_tsvector('english', coalesce(name, '')), 'A') ||
7           setweight(to_tsvector('english', coalesce(description, '')), 'B')
8         ) STORED;
9     SQL
10  end
11
12  def down
13    remove_column :companies, :searchable
14  end
15 end
```

Now let's add an index to this new `searchable` column of type `GIN`, doing it concurrently to avoid table locking:

RUBY

```
1 class AddIndexToCompaniesS < ActiveRecord::Migration[6.0]
2   disable_ddl_transaction!
3
4   def change
5     add_index :companies, :searchable, using: :gin, algorithm: :concurrently
6   end
7 end
```

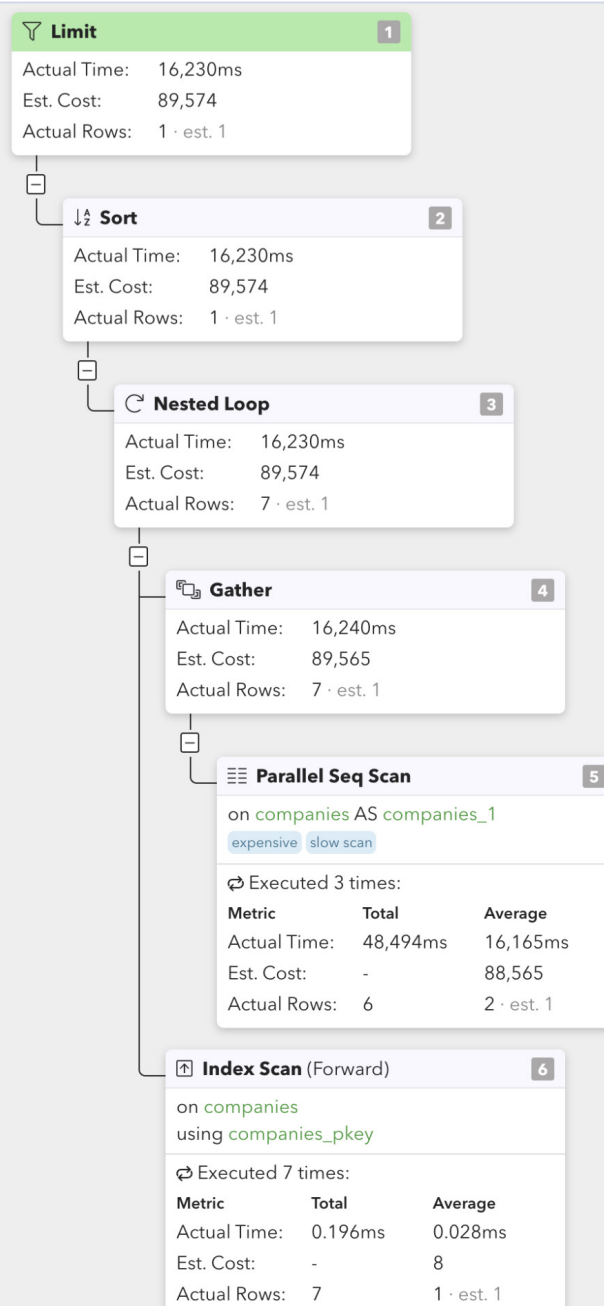
With our generated, stored column that is GIN indexed in place, we have to adjust the `pg_search` configuration to take **advantage of the tsvector column**:

RUBY

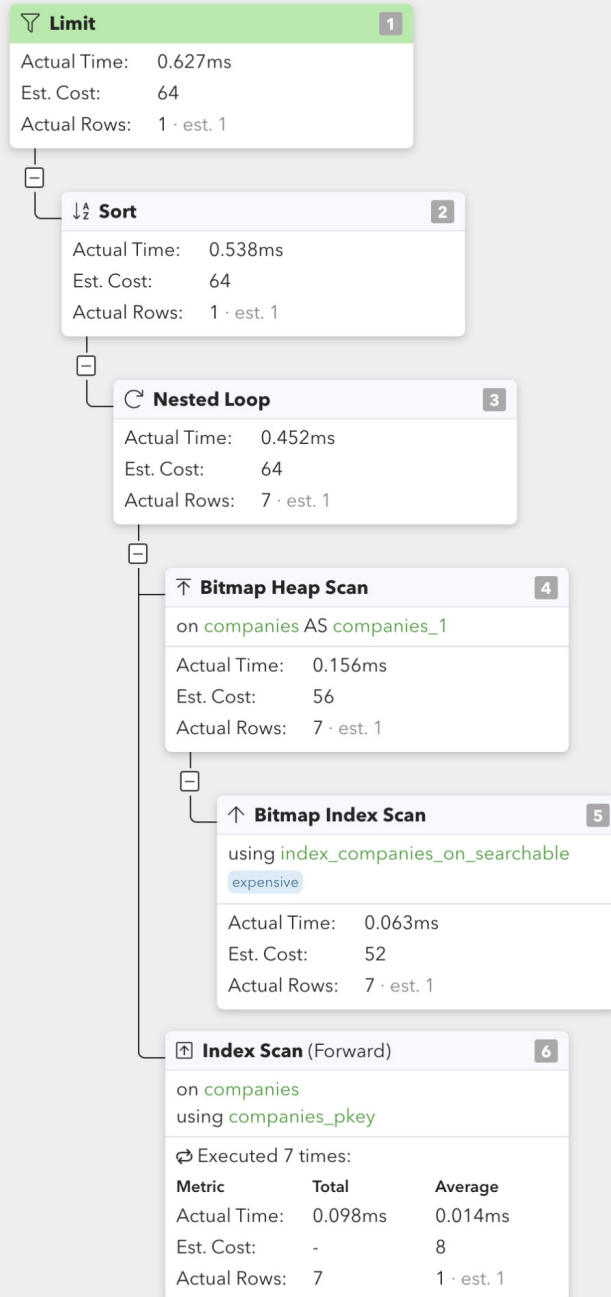
```
1 class Company < ApplicationRecord
2   include PgSearch::Model
3
4   pg_search_scope :search,
5     against: { name: 'A', description: 'B' },
6     using: {
7       tsearch: {
8         dictionary: 'english', tsvector_column: 'searchable'
9       }
10    }
11 end
```

To finish, let's look at the results. When we compare the EXPLAIN plans, we can see that we've gone from **16 seconds** to under **1 ms**, approximately **16000 times faster**. Talk about a speed improvement!

Before:



After:



Conclusion

With Postgres and Rails, we have the tools to get search right. Starting with exact matches, going to similarity searching using Trigrams, and ending up with Full Text Search allowing for natural language querying. For each of these options we have ways to optimize our queries by adding the appropriate indexes, giving us accurate and performant search results.

Try pganalyze for free

Have performance issues with your database, and looking for a way to improve indexes and query plans?

pganalyze provides deep insights for your database, with detailed performance analysis, and automated EXPLAIN plans, to quickly understand what is slow with your database. pganalyze integrates directly with major cloud providers, as well as self-managed Postgres installations.

Get started easily with a [free 14-day trial](#), or [learn more about our Enterprise product](#).

If you want, you can also [request a personal demo](#).



"Our overall usage of Postgres is growing, as is the amount of data we're storing and the number of users that interact with our products. pganalyze is essential to making our Postgres databases run faster, and makes sure end-users have the best experience possible."

Robin Fernandes, Software Development Manager
Atlassian

About pganalyze.

DBAs and developers use pganalyze to identify the root cause of performance issues, optimize queries and to get alerts about critical issues.

Our rich feature set lets you optimize your database performance, discover root causes for critical issues, get alerted about problems before they become big, gives you answers and lets you plan ahead.

Hundreds of companies monitor their production PostgreSQL databases with pganalyze.

Be one of them.

Sign up for a free trial today!

